

Module 3 Reflection

Grad Café data is self-reported and anonymous, which means the dataset reflects who chose to post rather than a representative sample of everyone who has applied to graduate school. Based on intuition, accepted applicants would be more likely to post than rejected applicants. Among rejected applicants, those with high scores, GPA, or other accolades would be more likely to post than those without the above. Rejections with high accolades might be more likely to post to show their accolades as a consolation. All of the aforementioned applicants would be more likely to post than an applicant who is waitlisted or needs to interview, since those statuses aren't "complete". Applicants who never discover the forum or do not feel motivated to document their process are absent from the data, and there is no way to measure how large this missing group is. Representation likely also varies by program and school, since active applicants are drawn disproportionately from whichever fields and audiences use it most. Self-reporting also introduces plain data-entry problems. Some examples of errors may stem from misremembering GPA or GRE scores or always rounding those figures up. The LLM portion of scraping seeks to standardize variations in naming schools or programs, but little can be done about an applicant's scores except ensure they exist within valid bounds.

These limitations matter when it comes to interpreting specific results. In this dataset, 79.26% of all recorded acceptances come from the Fall 2026 term alone, with every other term together accounting for barely a fifth of the total. That does not mean Fall 2026 was an unusually large or unusually successful admissions cycle compared to prior years. Instead, it reflects when people are actually active on Grad Café and which posts are seen. Applicants probably post while in their own cycle, checking the forum for other results and posting their own. The site shows the most recent results first, so scraping consecutive pages causes recency bias. The database can tell us that Fall 2026 acceptances make up the overwhelming majority of the acceptances recorded in the scraped dataset. However, this is not proof that graduate admissions activity is concentrated heavily in Fall 2026 in the real world. That distinction, between what a query returns and what real-world claims figures can support, is the central limitation of drawing conclusions from any self-selected, self-reported dataset.